

Imports

```
!pip install transformers

Collecting transformers
  Downloading transformers-4.12.5-py3-none-any.whl (3.1 MB)
    |████████████████████| 3.1 MB 7.1 MB/s
Collecting pyyaml>=5.1
  Downloading PyYAML-6.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2014_x86_64.whl (596 kB)
    |████████████████████| 596 kB 69.6 MB/s
Requirement already satisfied: filelock in /usr/local/lib/python3.7/dist-packages (from transformers) (3.4.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.7/dist-packages (from transformers) (1.21.0)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.7/dist-packages (from transformers) (4.62.3)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.7/dist-packages (from transformers) (2022.10.31)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from transformers) (2.28.1)
Collecting tokenizers<0.11,>=0.10.1
  Downloading tokenizers-0.10.3-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2014_x86_64.whl (3.3 MB)
    |████████████████████| 3.3 MB 49.8 MB/s
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.7/dist-packages (from transformers) (21.3)
Collecting huggingface-hub<1.0,>=0.1.0
  Downloading huggingface_hub-0.2.1-py3-none-any.whl (61 kB)
    |████████████████████| 61 kB 306 kB/s
Requirement already satisfied: importlib-metadata in /usr/local/lib/python3.7/dist-packages (from transformers) (4.2.0)
Collecting sacremoses
  Downloading sacremoses-0.0.46-py3-none-any.whl (895 kB)
    |████████████████████| 895 kB 75.8 MB/s
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.7/dist-packages (from huggingface-hub) (4.1.1)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging) (3.0.9)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata) (3.6.0)
Requirement already satisfied: urllib3!=1.25.0,!<1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests) (1.26.12)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests) (3.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests) (2022.9.24)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests) (3.0.4)
Requirement already satisfied: click in /usr/local/lib/python3.7/dist-packages (from sacremoses) (8.0.4)
Requirement already satisfied: joblib in /usr/local/lib/python3.7/dist-packages (from sacremoses) (1.1.0)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from sacremoses) (1.16.0)
Installing collected packages: pyyaml, tokenizers, sacremoses, huggingface-hub, transformers
  Attempting uninstall: pyyaml
    Found existing installation: PyYAML 3.13
    Uninstalling PyYAML-3.13:
      Successfully uninstalled PyYAML-3.13
Successfully installed huggingface-hub-0.2.1 pyyaml-6.0 sacremoses-0.0.46 tokenizers-0.10.3 transformers-4.12.5
```

```
# from preprocessing import preprocess
from transformers import DistilBertTokenizer, DistilBertModel
import torch
from torch import nn, optim
import copy
import random
import sklearn.metrics
import tqdm
import pickle
import pandas as pd
import numpy as np
import math
```

Helper Functions

```
def shuffle_data(input_1, input_2, labels):
    shuffled_input_1 = []
    shuffled_input_2 = []
    shuffled_labels = []
    indices = list(range(len(input_1)))
    random.shuffle(indices)
    for i in indices:
        shuffled_input_1.append(input_1[i])
        shuffled_input_2.append(input_2[i])
        shuffled_labels.append(labels[i])
    return (shuffled_input_1, shuffled_input_2, shuffled_labels)

def train(model, tokenizer, X_1, X_2, Y, learning_rate=0.1, batch_size=8, num_epochs=5):
    optimizer = optim.Adam(model.parameters(), lr=learning_rate)
    for epoch in range(num_epochs):
        total_loss = 0.0
        (shuffled_input_1, shuffled_input_2, shuffled_labels) = shuffle_data(X_1, X_2, Y)
```

```
for batch in tqdm.notebook.tqdm(range(0, len(X_1), batch_size), leave=False):
    #Randomly shuffle examples in each epoch
    input_1 = shuffled_input_1[batch:(batch + batch_size)]
    input_2 = shuffled_input_2[batch:(batch + batch_size)]
    encoded_input_1 = tokenizer(input_1, return_tensors='pt', is_split_into_words=False, padding=True)
    encoded_input_2 = tokenizer(input_2, return_tensors='pt', is_split_into_words=False, padding=True)
    labels = shuffled_labels[batch:(batch + batch_size)]
    labels_onehot = torch.zeros(len(labels), num_classes).cuda()
    for i in range(len(labels)):
        labels_onehot[i][labels[i]] = 1.0
    model.zero_grad()
    log_probs = model.forward(encoded_input_1, encoded_input_2, train=True)
    # print(log_probs)
    loss_batch = 0
    for idx in range(labels_onehot.shape[0]):
        loss_iteration = torch.neg(log_probs[idx]).dot(labels_onehot[idx])
        loss_batch += loss_iteration
    loss_batch /= labels_onehot.shape[0]
    loss_batch.backward()
    nn.utils.clip_grad_norm_(model.parameters(), 1.0)
    optimizer.step()
    total_loss += loss_batch.detach()
num_batches = math.ceil(len(X_1) / batch_size)
print(f"avg loss on epoch {epoch} = {total_loss / num_batches}")
```

```
def get_predictions(model, X_1, X_2, batch_size=8):
    all_predictions = np.array([])
    for batch in tqdm.notebook.tqdm(range(0, len(X_1), batch_size), leave=False):
        encoded_input_1 = tokenizer(X_1[batch:batch + batch_size], return_tensors='pt', is_split_into_words=False, padding=True)
        encoded_input_2 = tokenizer(X_2[batch:batch + batch_size], return_tensors='pt', is_split_into_words=False, padding=True)
        log_probs = model.forward(encoded_input_1, encoded_input_2, train=False)
        prediction_batch = torch.argmax(log_probs, dim=1)
        all_predictions = np.concatenate((all_predictions, prediction_batch.cpu().numpy()))
    return all_predictions
```

```
def get_predictions_cosine_similarity(model, tokenizer, X_1, X_2, threshold=0.96, batch_size=8):
    all_predictions = np.array([])
    for batch in tqdm.notebook.tqdm(range(0, len(X_1), batch_size), leave=False):
        encoded_input_1 = tokenizer(X_1[batch:batch + batch_size], return_tensors='pt', is_split_into_words=False, padding=True)
        encoded_input_2 = tokenizer(X_2[batch:batch + batch_size], return_tensors='pt', is_split_into_words=False, padding=True)
        pooler_output_1 = model(encoded_input_1['input_ids'].cuda(), encoded_input_1['attention_mask'].cuda()).last_hidden_state
        pooler_output_2 = model(encoded_input_2['input_ids'].cuda(), encoded_input_2['attention_mask'].cuda()).last_hidden_state
        cos = nn.CosineSimilarity(dim=1, eps=1e-6)
        output = cos(pooler_output_1, pooler_output_2).cpu().detach().numpy()
        preds = []
        for i in range(output.shape[0]):
            if output[i] > threshold:
                preds.append(1)
            else:
                preds.append(0)
        all_predictions = np.concatenate((all_predictions, preds))
    return all_predictions
```

```
def evaluate(Y, predictions):
    print("Accuracy: {}".format(sklearn.metrics.accuracy_score(Y, predictions)))
    print("F1 score: {}".format(sklearn.metrics.f1_score(Y, predictions)))
    print("Precision: {}".format(sklearn.metrics.precision_score(Y, predictions)))
    print("Recall: {}".format(sklearn.metrics.recall_score(Y, predictions)))
    print("Confusion matrix: \n{}\n".format(sklearn.metrics.confusion_matrix(Y, predictions)))
```

▼ Data Loader

```
# Raw csv
# df = pd.read_csv('train.csv')
# data = [list(df["question1"]), list(df["question2"]), list(df["is_duplicate"])]

# Preprocessed w/o transitivity
# with open('/content/drive/MyDrive/processed_data_wo_transitive.pkl', 'rb') as f:
#     data = pickle.load(f)

# Preprocessed w/ transitivity
# Finalized transitivity based on its superiority
with open('/content/drive/MyDrive/processed_data_2.pkl', 'rb') as f:
    data = pickle.load(f)

print("Original data has {} question pairs".format(len(data[0])))
```

Original data has 577240 question pairs

```
size = len(data[0])
dataset = data[:, :size]
train_ratio = 0.8
indices = list(range(size))
random.shuffle(indices)
train_indices = indices[:int(size*train_ratio)]
test_indices = indices[int(size*train_ratio):]
train_dataset = [[dataset[i][j] for j in train_indices] for i in range(len(dataset))]
test_dataset = [[dataset[i][j] for j in test_indices] for i in range(len(dataset))]

train_input_1 = [" ".join(train_dataset[0][i] for i in range(len(train_dataset[0])))
train_input_2 = [" ".join(train_dataset[1][i] for i in range(len(train_dataset[1])))
# train_input_1 = ["".join(train_dataset[0][i] for i in range(len(train_dataset[0])))
# train_input_2 = ["".join(train_dataset[1][i] for i in range(len(train_dataset[1])))
train_Y = train_dataset[2]
print(len(train_input_1), len(train_input_2), len(train_Y))
num_classes = 2

test_input_1 = [" ".join(test_dataset[0][i] for i in range(len(test_dataset[0])))
test_input_2 = [" ".join(test_dataset[1][i] for i in range(len(test_dataset[1])))
test_Y = test_dataset[2]
print(len(test_input_1), len(test_input_2), len(test_Y))
num_classes = 2

461792 461792 461792
115448 115448 115448
```

▼ Fine Tune Model

```
class SimilarityModelFineTuneBert(nn.Module):
    def __init__(self, dropout_rate=0.25):
        super(SimilarityModelFineTuneBert, self).__init__()
        self.bert = DistilBertModel.from_pretrained("distilbert-base-uncased").cuda()
        self.dropout = nn.Dropout(p=dropout_rate)
        self.feedforward_1 = nn.Linear(768*2, 300).cuda()
        self.non_lin_1 = nn.PReLU().cuda()
        self.dropout = nn.Dropout(p=dropout_rate)
        self.feedforward_2 = nn.Linear(300, 300).cuda()
        self.non_lin_2 = nn.PReLU().cuda()
        self.dropout = nn.Dropout(p=dropout_rate)
        self.feedforward_3 = nn.Linear(300, 2).cuda()
        self.log_softmax = nn.LogSoftmax(dim=0).cuda()

    def forward(self, encoded_input_1, encoded_input_2, train=False):
        if train:
            self.train()
        else:
            self.eval()
        pooler_output_1 = self.bert(encoded_input_1['input_ids'].cuda(), encoded_input_1['attention_mask'].cuda())
        pooler_output_2 = self.bert(encoded_input_2['input_ids'].cuda(), encoded_input_2['attention_mask'].cuda())
        # print(pooler_output_1)
        concatenated_output = torch.cat([pooler_output_1, pooler_output_2], axis=1).cuda()
        f1 = self.dropout(self.non_lin_1(self.feedforward_1(concatenated_output)))
        f2 = self.dropout(self.non_lin_2(self.feedforward_2(f1)))
        return self.log_softmax(self.feedforward_3(f2))
```

▼ Config

```
dropout_rate = 0.25
batch_size = 128
learning_rate = 0.0001
num_epochs = 5
```

▼ Train

```
tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')
print("Training fine tune model")
fine_tune_model = SimilarityModelFineTuneBert()
train(fine_tune_model, tokenizer, train_input_1, train_input_2, train_Y, learning_rate=learning_rate, num_epochs=r
```

```

    Downloading: 100%                226k/226k [00:00<00:00, 941kB/s]

    Downloading: 100%                28.0/28.0 [00:00<00:00, 668B/s]

    Downloading: 100%                455k/455k [00:00<00:00, 878kB/s]

    Downloading: 100%                483/483 [00:00<00:00, 15.9kB/s]
Training fine tune model
    Downloading:                    256M/256M [00:05<00:00,
100%                               47.5MB/s]
Some weights of the model checkpoint at distilbert-base-uncased were not us
- This IS expected if you are initializing DistilBertModel from the checkpc
- This IS NOT expected if you are initializing DistilBertModel from the che
100%                               3608/3608 [21:36<00:00, 3.20it/s]
avg loss on epoch 0 = 4.696255683898926
100%                               3608/3608 [21:37<00:00, 3.00it/s]
avg loss on epoch 1 = 4.579102039337158
100%                               3608/3608 [21:37<00:00, 3.23it/s]
avg loss on epoch 2 = 4.508677005767822
100%                               3608/3608 [21:36<00:00, 2.84it/s]
avg loss on epoch 3 = 4.456218719482422
.....
with open("bert-ffn-fine-tune-pickle.pkl", 'wb') as f:
    pickle.dump(fine_tune_model, f)

torch.save(fine_tune_model.state_dict(), 'bert-ffn-fine-tune-state-dict.pt')
```

▼ Eval

```

print("Evaluating fine tune model on train dataset")
predictions = get_predictions(fine_tune_model, train_input_1, train_input_2, batch_size=batch_size)
evaluate(train_Y, predictions)

print("Evaluating fine tune bert model on test dataset")
predictions = get_predictions(fine_tune_model, test_input_1, test_input_2, batch_size=batch_size)
evaluate(test_Y, predictions)

    Evaluating fine tune model on train dataset
    100%                3608/3608 [10:01<00:00, 6.66it/s]
    Accuracy: 0.9403844154944218
    F1 score: 0.9187296676565923
    Precision: 0.8887264336830563
    Recall: 0.9508294888637683
    Confusion matrix:
    [[278654  19483]
     [  8047 155608]]

    Evaluating fine tune bert model on test dataset
    100%                902/902 [02:30<00:00, 5.74it/s]
    Accuracy: 0.8613141847411822
    F1 score: 0.8138190867123271
    Precision: 0.7766040080782973
    Recall: 0.8547803996287068
    Confusion matrix:
    [[64444 10066]
     [ 5945 34993]]
```

▼ Cosine Similarity Baseline (Discontinued)

```
# Stopped working on this after evaluating a baseline

tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')
vanilla = DistilBertModel.from_pretrained("distilbert-base-uncased").cuda()
print("Evaluating vanilla model on train dataset")
print("NOTE: For vanilla model with cosine based similarity, train-test split doesn't matter\n")
predictions = get_predictions_cosine_similarity(vanilla, tokenizer, train_input_1, train_input_2, 0.96)
evaluate(train_Y, predictions)

print("Evaluating vanilla model on test dataset")
print("NOTE: For vanilla model with cosine based similarity, train-test split doesn't matter\n")
predictions = get_predictions_cosine_similarity(vanilla, tokenizer, test_input_1, test_input_2, 0.96)
evaluate(test_Y, predictions)
```

Some weights of the model checkpoint at distilbert-base-uncased were not us  
- This IS expected if you are initializing DistilBertModel from the checkpc  
- This IS NOT expected if you are initializing DistilBertModel from the che  
Evaluating vanilla model on train dataset  
NOTE: For vanilla model with cosine based similarity, train-test split does

99%994/1000 [00:15<00:00, 62.20it/s]

Accuracy: 0.598375  
F1 score: 0.6064911206368647  
Precision: 0.4746932515337423  
Recall: 0.8396066463207867  
Confusion matrix:  
[[2311 2740]  
 [ 473 2476]]

Evaluating vanilla model on test dataset  
NOTE: For vanilla model with cosine based similarity, train-test split does

100%249/250 [00:03<00:00, 61.54it/s]

Accuracy: 0.6035  
F1 score: 0.6170931916948333  
Precision: 0.4881588999236058  
Recall: 0.8385826771653543  
Confusion matrix:  
[[568 670]  
 [123 639]]

### ▼ Static BERT Model (Discontinued)

# Stopped training this after superiority of finetuning BERT was realized

```
class SimilarityModelStaticBert(nn.Module):
    def __init__(self):
        super(SimilarityModelStaticBert, self).__init__()
        self.bert = DistilBertModel.from_pretrained("distilbert-base-uncased").cuda()
        for param in self.bert.parameters():
            param.requires_grad = False
        self.feedforward_1 = nn.Linear(768*2, 300).cuda()
        self.non_lin_1 = nn.PReLU().cuda()
        self.feedforward_2 = nn.Linear(300, 300).cuda()
        self.non_lin_2 = nn.PReLU().cuda()
        self.feedforward_3 = nn.Linear(300, 2).cuda()
        self.log_softmax = nn.LogSoftmax(dim=0).cuda()

    def forward(self, encoded_input_1, encoded_input_2, train=False):
        if train:
            self.train()
        else:
            self.eval()
        pooler_output_1 = self.bert(encoded_input_1['input_ids'].cuda(), encoded_input_1['attention_mask'].cuda())
        pooler_output_2 = self.bert(encoded_input_2['input_ids'].cuda(), encoded_input_2['attention_mask'].cuda())
        # print(pooler_output_1)
        concatenated_output = torch.cat([pooler_output_1, pooler_output_2], axis=1).cuda()
        f1 = self.non_lin_1(self.feedforward_1(concatenated_output))
        f2 = self.non_lin_2(self.feedforward_2(f1))
        return self.log_softmax(self.feedforward_3(f2))
```

```
tokenizer = DistilBertTokenizer.from_pretrained('distilbert-base-uncased')
print("Training static bert model")
static_bert_model = SimilarityModelStaticBert()
train(static_bert_model, tokenizer, train_input_1, train_input_2, train_Y, learning_rate=0.0001, num_epochs=10, ba
```

```
Training static bert model
Some weights of the model checkpoint at distilbert-base-uncased were not us
- This IS expected if you are initializing DistilBertModel from the checkpc
- This IS NOT expected if you are initializing DistilBertModel from the che

99% 248/250 [00:11<00:00, 21.42it/s]
avg loss on epoch 0 = 0.10724116116762161
100% 249/250 [00:11<00:00, 21.18it/s]
avg loss on epoch 1 = 0.10655724257230759
100% 249/250 [00:11<00:00, 21.37it/s]
avg loss on epoch 2 = 0.10627871006727219
99% 248/250 [00:11<00:00, 21.48it/s]
avg loss on epoch 3 = 0.10622921586036682
99% 248/250 [00:11<00:00, 22.28it/s]
avg loss on epoch 4 = 0.10600847005844116

print("Evaluating static bert model on train dataset")
predictions = get_predictions(static_bert_model, train_input_1, train_input_2, batch_size=32)
evaluate(train_Y, predictions)

print("Evaluating static bert model on test dataset")
predictions = get_predictions(static_bert_model, test_input_1, test_input_2, batch_size=32)
evaluate(test_Y, predictions)

Evaluating static bert model on train dataset
100% 250/250 [00:10<00:00, 23.26it/s]
Accuracy: 0.70625
F1 score: 0.6392385630948726
Precision: 0.5840112201963534
Recall: 0.7060020345879959
Confusion matrix:
[[3568 1483]
 [ 867 2082]]

Evaluating static bert model on test dataset
98% 62/63 [00:02<00:00, 23.51it/s]
Accuracy: 0.6835
F1 score: 0.6165960024227742
Precision: 0.5725534308211474
Recall: 0.6679790026246719
Confusion matrix:
[[858 380]
 [253 509]]
```